


Best Practices for Writing SQL in PL/SQL



Steven Feuerstein

PL/SQL Evangelist, Quest Software

www.ToadWorld.com/SF

www.plsqlchallenge.com

steven.feuerstein@quest.com

How to benefit most from this session

- Watch, listen, *ask questions*. Then afterwards....
- Download and use any of my the training materials, available at my "cyber home" on Toad World, a portal for Toad Users and PL/SQL developers:

PL/SQL Obsession

<http://www.ToadWorld.com/SF>

- Download and use any of my scripts (examples, performance scripts, reusable code) from the demo.zip, available from the same place.

filename_from_demo_zip.sql
- You have my permission to use *all* these materials to do internal trainings and build your own applications.
 - But they should not considered production ready.
 - You must test them and modify them to fit your needs.

And some other incredibly fantastic and entertaining websites for PL/SQL



Best Practices for Writing SQL in PL/SQL

- **Set standards and guidelines for writing SQL.**
- **Take full advantage of the SQL language.**
- **Hide SQL statements behind an interface.**
- **Hide all tables in schemas users cannot access.**
- **Qualify every identifier in SQL statements.**
- **Dot-qualify references to Oracle-supplied objects.**
- **Use SELECT INTO for single row fetches.**
- **Always BULK COLLECT with LIMIT clause.**
- **Always use FORALL for multi-row DML.**
- **Use collection and TABLE operator for IN clauses of indeterminate count.**
- **Avoid implicit conversions.**
- **Key dynamic SQL best practices**



Set standards and guidelines for SQL

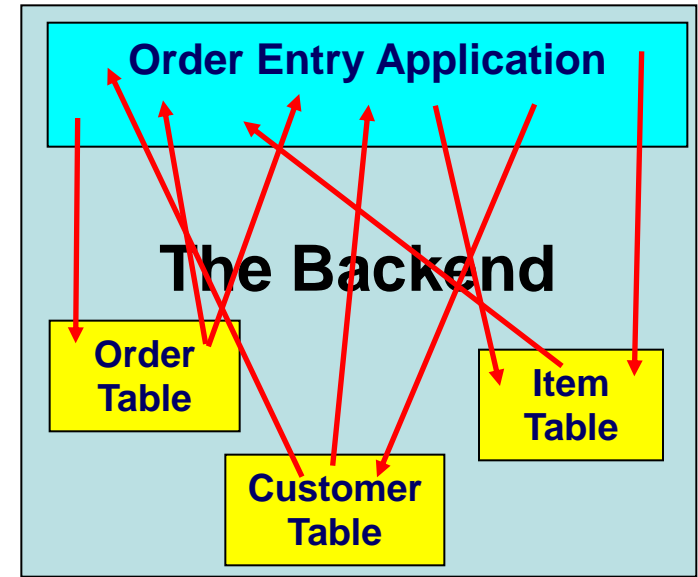
- **Many organizations have coding standards.**
 - How to format code, how to name programs and variables, etc.
- **Very few development teams have standards for how, when and where to write SQL statements.**
 - We just all take SQL for granted.
- **This is very strange and very dangerous.**

Why lack of standards for SQL is dangerous

- **SQL statements are among the most critical elements of our applications.**
- **SQL statements reflect our business model.**
 - And those models are always changing.
- **SQL statements cause most of the performance problems in our applications.**
 - Tuning SQL and the way that SQL is called in PL/SQL overwhelms all other considerations.
- **Many runtime errors result from integrity and check constraints on tables.**

The fundamental problem with SQL in PL/SQL

- **We take it entirely for granted.**
 - Why not? It's so easy to write SQL in PL/SQL!
- **As a result, our application code is packed full of SQL statements, with many repetitions and variations.**
 - Worst of all: SQL in .Net and Java!
- **This makes it very difficult to optimize and maintain the application code.**



So set some SQL standards!

- **At a minimum, before starting next application, ask yourselves explicitly:**
 - Do we want standards or should we just do whatever we want, whenever we want?
 - That way, you are making a conscious decision.
- **This presentation as a whole forms a reasonable foundation for such standards.**
- **Another excellent resource (and source for this presentation):**

Doing SQL from PL/SQL: Best and Worst Practices
by Bryn Llewellyn, PL/SQL Product Manager

http://www.oracle.com/technology/tech/pl_sql/pdf/doing_sql_from_plsql.pdf

Fully leverage SQL in your PL/SQL code

- Oracle continually adds significant new functionality to the SQL language.
- If you don't keep up with SQL capabilities, you will write slower, more complicated PL/SQL code than is necessary.
 - I am actually a good example of what you *don't* want to do or how to be.
- So take the time to refresh your understanding of Oracle SQL in 10g and 11g.

Some exciting recently added SQL features

- **Courtesy of Lucas Jellama of AMIS Consulting**
- **Analytical Functions**
 - Primarily LAG and LEAD; these allow to look to previous and following rows to calculate differences)
- **WITH clause (subquery factoring)**
 - Allows the definition of 'views' inside a query that can be used and reused; they allow procedural top-down logic inside a query
- **Flashback query**
 - No more need for journal tables, history tables, etc.
- **ANSI JOIN syntax**
 - Replaces the (+) operator and introduces FULL OUTER JOIN
- **SYS_CONNECT_BY_PATH and CONNECT_BY_ROOT for hierarchical queries**
- **Scalar subquery**
 - Adds a subquery to a query like a function call.

```
select d.deptno
       , (select count(*)
          from emp e where
          e.deptno = d.deptno)
       number_staff from dept
```

Hide SQL statements behind an interface

- You, of course, need to write SQL in Oracle applications.
- And PL/SQL is the best place to write and store the SQL statements.
- But we must stop writing SQL statements *all over* the application code base.
 - Repetition of SQL is a real nightmare.
- The best way to understand this is to accept a harsh reality:


Every SQL statement you write is a hard-coding that is *worse* than a hard-coded literal.

SQL as Hard-Coding....huh?

- **We all agree that hard-coding "magic values" is a bad idea.**
 - When value changes (and it will), you must find all occurrences and update them.
- **But SQL statements suffer from the same problem.**
 - When you write SQL, you are saying "Today, *at this moment*, this is the complex code needed to describe this dataset."
- **There is no logical distinction between a magic value and a "magic query."**
 - Both will change, both should not be repeated.

SQL as a Service

- Think of SQL as a *service* that is provided to you, not something you write.
 - Or if you write it, you put it somewhere so that it can be easily found, reused, and maintained.
- This service consists of *programs* defined in the data access layer.
 - Known as table APIs, transaction APIs, or data encapsulation, these programs contain all the intelligence about business transactions and underlying tables.



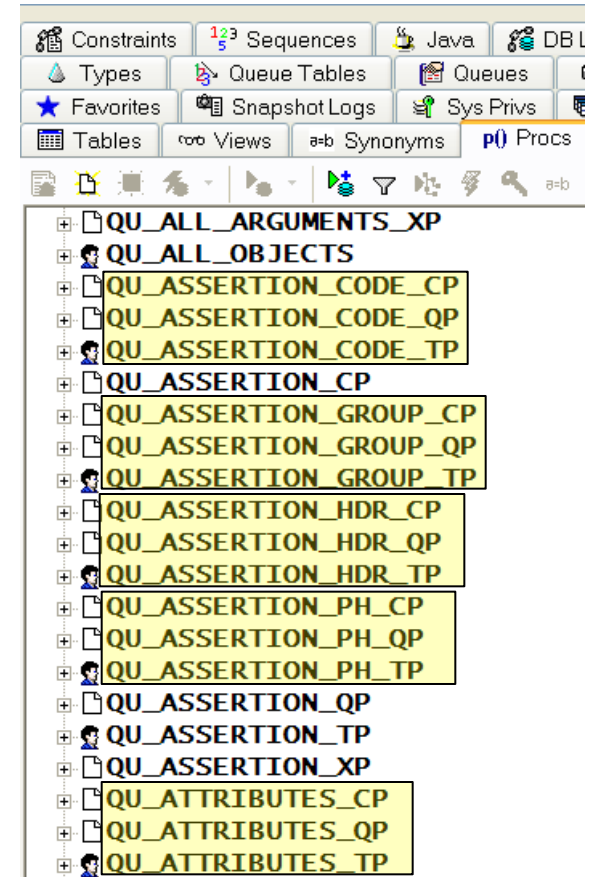
With encapsulated SQL I can...

- **Change/improve my implementation with minimal impact on my application code.**
 - The underlying data model is constantly changing.
 - We can depend on Oracle to add new features.
 - We learn new ways to take advantage of PL/SQL.
- **Vastly improve my SQL-related error handling.**
 - Do you handle `dup_val_on_index` for INSERTs, `too_many_rows` for SELECT INTOs, etc?
- **Greatly increase my productivity**
 - I want to spend as much time as possible implementing business requirements.

11g_emplu.pkg

Example: Quest Code Tester backend

- For each table, we have three generated packages:
 - <table>_CP for DML
 - <table>_QP for queries
 - <table>_TP for types
- And usually an "extra stuff" package (_XP) with custom SQL logic and related code.
 - You can't generate everything.



How to implement data encapsulation

(After all, I did promise to be practical!)

- It must be very consistent, well-designed and efficient - or it will not be used.
- **Best solution: generate as much of the code as possible.**
 - And any custom SQL statements should be written once and placed in a standard *container* (usually a package).
- **One option for generating table APIs for use with PL/SQL is the freeware Quest CodeGen Utility, available at PL/SQL Obsession:**

www.ToadWorld.com/SF

Encapsulating Data Retrieval

- **Simplify query access with views**
 - Another form of encapsulation
- **Hide queries behind functions**
 - Return cursor variable to non-PL/SQL host environment.
 - Return collection or record to other PL/SQL programs
- **Use table functions to encapsulate complex data transformations**

Encapsulating DML statements

- **Hiding inserts are relatively straightforward**
 - Insert by record, collection, individual columns
- **Encapsulating updates is more challenging.**
 - Many variations
 - Some choose a blend of dynamic and static SQL
 - Others use a parallel "indicator" argument to specify which of the columns should be included in the update.
 - You *will* write your own custom encapsulators.

Before you start your next application...

- **Sit down as a team and decide what you are going to do about SQL.**

Choice #1. Keep doing what you've been doing (everyone writes SQL wherever and whenever they want.

Choice #2. Full encapsulation: no directly granted privileges on tables, only access path is through API.

Choice #3. Encapsulate most important tables and run validations against code to identify violations of the API.

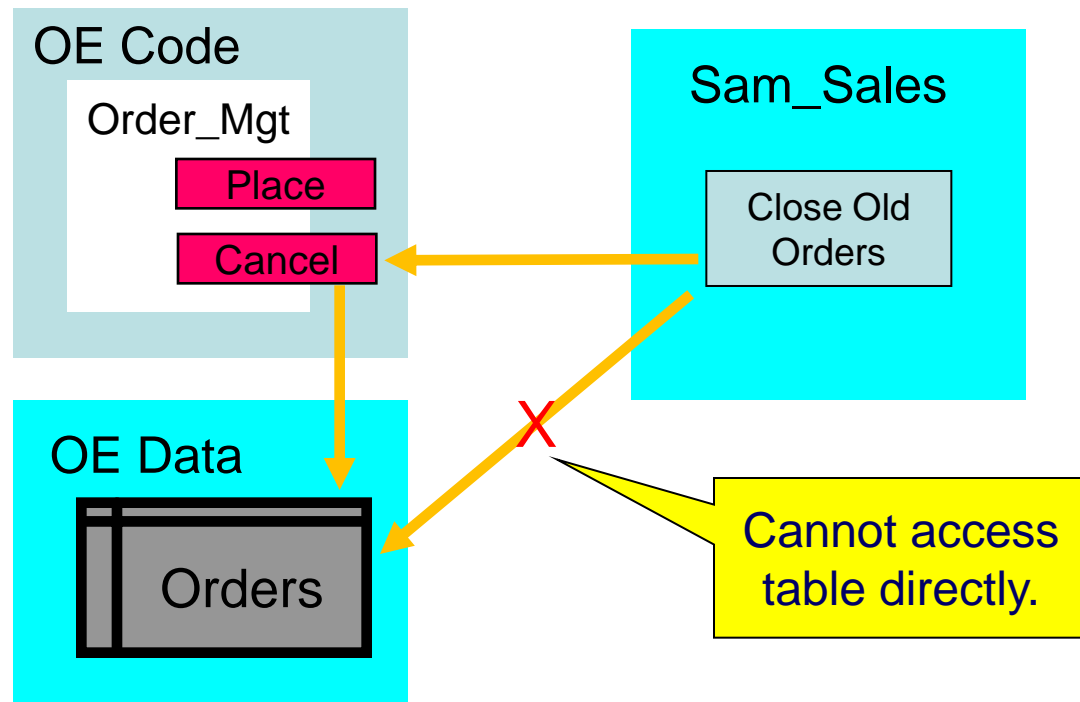
Choice #4. Encapsulate queries to prepare for upgrade to Oracle11g and the function result cache.

Hide all tables in schemas users cannot access.

- **A fundamental issue of control and security.**
- **Do not allow users to connect to any schema that contains tables.**
 - Simply too risky.
- **Define tables in other schemas.**
- **Grant access via privileges, mostly via EXECUTE on packages that maintain the tables.**

Architecture with inaccessible schema for data

- The OE Data schemas own all tables.
- The OE Code schema owns the code and has directly granted privileges on the tables.
- User schemas have execute authority granted on the code.



Qualify every column and identifier in the SQL statement.

- Improves readability.
- Avoids potential bugs when variable names match column names.
- Minimizes invalidation of dependent program units in Oracle11g.

11g_fgd*.sql

Instead of this....

```
PROCEDURE abc (...)  
IS  
BEGIN  
    SELECT last_name  
        INTO l_name  
        FROM employees  
        WHERE employee_id = employee_id_in;
```

Write this....

```
PROCEDURE abc (...)  
IS  
BEGIN  
    SELECT e.last_name  
        INTO l_name  
        FROM employees e  
        WHERE e.employee_id = abc.emp_id_in;
```

Dot-qualify all references to Oracle-supplied objects with "SYS."

- Another annoying, but incontestable recommendation.
- If you don't prefix calls to all supplied packages with "SYS.", you are more vulnerable to *injection*.

```
BEGIN
  run_dynamic_plsql_block
    (append_this_in =>
      'employee id=101; EXECUTE IMMEDIATE
        ''CREATE OR REPLACE PACKAGE DBMS_OUTPUT ... ''
    );
END;
```

Use **SELECT INTO** for single row fetches.

- Long ago, Oracle "gurus" warned against **SELECT INTO** (implicit query) and pushed explicit cursors for single row fetches.
- Then Oracle optimized **SELECT INTO**, so that implicits are generally faster than explicits.
 - In Oracle11, the difference is small.
- The most important thing to do, however, is to *hide* your query inside a function.
 - So when Oracle changes the picture again, you only have to adjust your code in one place (for each query).

Always BULK COLLECT with LIMIT clause.

- **First, always use BULK COLLECT to retrieve multiple rows of data.**
 - Note: "Read only" (no DML) cursor FOR loops are automatically optimized to array performance.
- **With BULK COLLECT into a varray if you *know* there is a maximum limit on the number of rows retrieved.**
- **Otherwise, use the LIMIT clause with your BULK COLLECT statement.**
 - Avoid hard-coding: LIMIT can be a variable or parameter.

Always use **FORALL** for multi-row DML.

- **Convert all loops containing DML statements into **FORALL** statements.**
 - Incredible boost in performance.
- **The conversion process can be tricky and complicated.**
 - Use **SAVE EXCEPTIONS** or **LOG ERRORS** to continue past errors.
 - Use **INDICES OF** and **VALUES OF** with sparse collections

Use collection and TABLE operator for IN clauses of indeterminate count.

- **The IN clause may contain no more than 1000 elements.**
- **Several options for "dynamic" IN clause:**
 - Dynamic SQL
 - IN clause with TABLE operator
 - MEMBER OF
- **Using a collection with the TABLE operator offers best flexibility and performance.**
 - Must be declared at schema level.

Avoid implicit conversions

- **Oracle is very forgiving.**
 - If it can implicitly convert a value from one datatype to another, it will do it without complaint.
- **There is, however, a price to pay.**
 - Implicit conversions can affect optimization of SQL statements.
 - There is overhead to the conversion that is best avoided.
- **So whenever possible....**
 - Use correct datatypes to avoid need to convert.
 - Rely on explicit rather than implicit conversions.

Key Dynamic SQL Best Practices

- **Always EXECUTE IMMEDIATE a variable.**
 - Otherwise it will be very difficult to debug those complicated strings.
- **Stored programs with dynamic SQL should be AUTHID CURRENT_USER.**
 - Make sure the right DB objects are affected.
- **Dynamic DDL programs should be autonomous transactions.**
 - Watch out for those implicit commits!
- **Minimize the possibility of SQL injection.**

dropwhatever.sp

SQL (code) Injection

- **"Injection" means that unintended and often malicious code is inserted into a dynamic SQL statement.**
 - Biggest risk occurs with dynamic PL/SQL, but it is also possible to subvert SQL statements.
- **Best ways to avoid injection:**
 - Restrict privileges tightly on user schemas.
 - Use bind variables whenever possible.
 - Check dynamic text for dangerous text.
 - Use DBMS_ASSERT to validate object names, like tables and views.

code_injection.sql
sql_guard.*
dbms_assert_demo.sql

usebinding.sp
toomuchbinding.sp
useconcat*.*
ultrabind.*



Best Practices for Writing SQL in PL/SQL

- **Stop taking SQL for granted.**
 - The most important part of your application code base.
- **Fully utilize the SQL language.**
 - If you can do it in SQL, don't complicate with PL/SQL.
- **Avoid repetition of SQL statements.**
 - At a minimum, hide queries inside functions to prepare for the function result cache.
- **Take advantage of collections for flexibility and performance .**
 - FORALL and BULK COLLECT
 - IN clause flexibility